

HW 5

This assignment covers Comparison of Decision Trees and Support Vector Machine. **DO NOT ERASE MARKDOWN CELLS AND INSTRUCTIONS IN YOUR HW submission**

- **Q** - QUESTION
- **A** - Where to input your answer

Instructions

Keep the following in mind for all notebooks you develop:

- Structure your notebook.
- Use headings with meaningful levels in Markdown cells, and explain the questions each piece of code is to answer or the reason it is there.
- Make sure your notebook can always be rerun from top to bottom.
- Objective of this assignment is to help you master python and scikit-learn package.
- See [README.md](#) for homework submission instructions

Related Tutorials

- [Decision Tree with KFold Cross Validation](#)
- [Decision Tree with Bagging](#)
- [Support Vector Machine](#)

Data Processing

Q1 Get training data from the dataframe

1. Load HW5_data.csv from ``data" folder into the dataframe
2. Check if there is any NaN in the dataset
3. Remove the rows with NaN values.
4. Print how many examples belong to each class in the data frame.

A1 Replace ??? with code in the code cell below

```
In [ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split

#Read the data file using the appropriate separator as input to read_csv()

df = pd.read_csv('../data/HW5_data.csv', na_values='', sep=",")
df.head(10)
```

```
Out[ ]:
```

	Mean of the integrated profile	Standard deviation of the integrated profile	Excess kurtosis of the integrated profile	Skewness of the integrated profile	Mean of the DM- SNR curve	Standard deviation of the DM-SNR curve	Excess kurtosis of the DM-SNR curve	Skewness of the DM-SNR curve	target_class
0	121.156250	48.372971	0.375485	-0.013165	3.168896	18.399367	7.449874	65.159298	
1	76.968750	36.175557	0.712898	3.388719	2.399666	17.570997	9.414652	102.722975	
2	130.585938	53.229534	0.133408	-0.297242	2.743311	22.362553	8.508364	74.031324	
3	156.398438	48.865942	-0.215989	-0.171294	17.471572	NaN	2.958066	7.197842	
4	84.804688	36.117659	0.825013	3.274125	2.790134	20.618009	8.405008	76.291128	
5	121.007812	47.176944	0.229708	0.091336	2.036789	NaN	9.546051	112.131721	
6	79.343750	42.402174	1.063413	2.244377	141.641304	NaN	-0.700809	-1.200653	
7	109.406250	55.912521	0.565106	0.056247	2.797659	19.496527	9.443282	97.374578	
8	95.007812	40.219805	0.347578	1.153164	2.770067	18.217741	7.851205	70.801938	
9	109.156250	47.002234	0.394182	0.190296	4.578595	NaN	5.702532	36.342493	

```
In [ ]: # check if there is NaN in the dataset
df.isnull().sum()
```

```
Out[ ]: Mean of the integrated profile          0
Standard deviation of the integrated profile  0
Excess kurtosis of the integrated profile    1735
Skewness of the integrated profile          0
Mean of the DM-SNR curve                   0
Standard deviation of the DM-SNR curve     1178
Excess kurtosis of the DM-SNR curve        0
Skewness of the DM-SNR curve               625
target_class                              0
dtype: int64
```

```
In [ ]: #Drop NaNs if there is any
df.dropna(inplace=True)

# Count number of entries for different target_class
df.target_class.count()
```

```
Out[ ]: 9273
```

Q2 Separate training and testing data from the dataframe

1. Assign values of `target_class` column to `y`, note you have to use `.values` method
2. Drop `target_class` column from data frame,
3. Assign df values to `x`
4. Split dataset into train and test data use `train_test_split` with `test_size = 0.25`, stratify `y` and `random_state = 1238`

A2 Replace ??? with code in the code cell below

```
In [ ]: # Assign values of ``target_class`` column to y, note you have to use .values method
y = df.target_class.values
# Drop 'target_class' column from data frame,
df.drop(columns=['target_class'], inplace=True)
# Assign df values to x
x = df.values
# View shape of x and y
print('Shape of x = ', x.shape)
print('Shape of y = ', y.shape)

xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.25, random_state=1238)

Shape of x = (9273, 8)
Shape of y = (9273,)
```

```
In [ ]: xtrain.shape, xtest.shape, ytrain.shape, ytest.shape
```

```
Out[ ]: ((6954, 8), (2319, 8), (6954,), (2319,))
```

Decision Tree

Decision Tree with different depth

Q3 Train `DecisionTreeClassifier` Model at different depths

1. Create four `DecisionTreeClassifier` models with different parameters. Use `max_depth` size = 1, 2, 5, 25 & `max_leaf_nodes`=5, 10, 15, 25 respectively
2. Use `random_state=30` & `criterion='entropy'` for all models
3. Fit the four different models with the train data.
4. Predict the test data using trained models
5. Calculate the Mean Squared Error(MSE) of each model's prediction
6. Print precision recall curve for the test data with the minimum MSE value from four trained models.

A3 Replace ??? with code in the code cell below

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
        from sklearn.metrics import mean_squared_error

        #create decision tree classifier
        clf_1 = DecisionTreeClassifier(criterion="entropy", max_depth=1, max_leaf_nodes=5,
        clf_2 = DecisionTreeClassifier(criterion="entropy", max_depth=2, max_leaf_nodes=10,
        clf_3 = DecisionTreeClassifier(criterion="entropy", max_depth=5, max_leaf_nodes=15,
        clf_4 = DecisionTreeClassifier(criterion="entropy", max_depth=25, max_leaf_nodes=25

        #fit classifier model
        clf_1.fit(xtrain,ytrain)
        clf_2.fit(xtrain,ytrain)
        clf_3.fit(xtrain,ytrain)
        clf_4.fit(xtrain,ytrain)

        #predict
        pred_1 = clf_1.predict(xtest)
        pred_2 = clf_2.predict(xtest)
        pred_3 = clf_3.predict(xtest)
        pred_4 = clf_4.predict(xtest)

        #calculate mean_squared_error
        print('clf_1 MSE: ', mean_squared_error(ytest, pred_1))
        print('clf_2 MSE: ', mean_squared_error(ytest, pred_2))
        print('clf_3 MSE: ', mean_squared_error(ytest, pred_3))
        print('clf_4 MSE: ', mean_squared_error(ytest, pred_4))

        clf_1 MSE:  0.0258732212160414
        clf_2 MSE:  0.0258732212160414
        clf_3 MSE:  0.02501078050884002
        clf_4 MSE:  0.0258732212160414
```

Precision-Recall Curve for Best Above

Important Note: If `from_estimator()` function gives Attribute error then it means your sklearn is not updated.

- If you are using conda, you can upgrade with

conda upgrade -c conda-forge scikit-learn


- or, with pip,


python -m pip install scikit-learn --upgrade

```
In [ ]: # Use the below one
from sklearn.metrics import precision_recall_curve

# Or this below one, whichever suits you
from sklearn.metrics import PrecisionRecallDisplay
import matplotlib.pyplot as plt

# using prediction from clf_3 because it had the MSE closest to 0
disp = PrecisionRecallDisplay.from_estimator(clf_3, xtest, ytest)
disp.plot()
plt.show()
```





Decision Tree with K-fold cross validation

Q4 Use Kfold on the test dataset, and evaluate the best model

1. Use `cross_val_score` and fit your best model with `k = 5` fold size on test data
2. Calculate average scores in kfold

A4 Replace ??? with code in the code cell below

```
In [ ]: from sklearn.model_selection import KFold, cross_val_score
```

```
scores = cross_val_score(clf_4, xtest, ytest, cv=5)
print("Cross-validation scores: {}".format(scores))
print("Average cross-validation score: {}".format(scores.sum()/5))
```

```
Cross-validation scores: [0.95689655 0.9762931 0.96982759 0.95258621 0.9524838 ]
Average cross-validation score: 0.9616174499143517
```

Decision Tree with Bagging

Q5 Now we will use Bagging technique on the our previous best model, and evaluate it

Part 1:

1. Now, Create a Bagged Model passing `model = previous_best`, `n_estimators = 10` & `random_state=1` to `BaggingClassifier()`
2. Fit the model with the train data
3. Predict the values with the test data
4. Calculate the test MSE
5. Plot Precision-Recall Curve from the true & predicted test data

A5 Replace ??? with code in the code cell below

```
In [ ]: from sklearn.ensemble import BaggingClassifier


# Use BaggingClassifier to fit the training data
# Calculate the mean squared error

#Load BaggingClassifier model and pass n_estimators=10, random_state=1
bagged_clf = BaggingClassifier(estimator=clf_3, n_estimators=10, random_state=1)
bagged_clf.fit(xtrain, ytrain)
pred = bagged_clf.predict(xtest)
print('bagged_clf MSE: ', mean_squared_error(ytest, pred))
```

bagged_clf MSE: 0.02630444156964209

```
In [ ]: #pass necessary parameters to PrecisionRecallDisplay.from_estimator()

PrecisionRecallDisplay.from_estimator(bagged_clf, xtest, ytest)
plt.show()
```



Part 2:

1. Why BaggingClassifier is called an ensemble technique? why it works better most of the time than the single model classifiers?

BaggingClassifier is called an ensemble technique because it combines the predictions of multiple models in making its final prediction. In the case of our `bagging_clf`, we incorporated a `DecisionTreeClassifier` into the "estimator" parameter of the `BaggingClassifier`. Ensemble techniques generally work better than single model classifiers because the combination of predictions from multiple models improves generalization as well as reducing variance and variability.

1. What is the disadvantage of increasing the number of estimators while using BaggingClassifier? Explain with an appropriate example.

Increasing the number of estimators can lead to potential overfitting, which is exactly what bagging is supposed to reduce the risk of.

Support Vector Machine(SVM)

Q6 Create SVM Model on the training set, and do the following

Part:1

1. Now, Create a SVM Model with default parameters
2. Fit the model with the train data
3. Predict the values with the test data
4. Calculate the model accuracy on test data
5. Plot confusion matrix on the test data (Make font size 16)

A6 Replace ??? with code in the code cell below

```
In [ ]: # import SVC classifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

svc = SVC()

# fit classifier to training set
svc.fit(xtrain, ytrain)

# make predictions on test set
svc_pred = svc.predict(xtest)


# compute and print accuracy score
acc = accuracy_score(ytest, svc_pred)
print('Model accuracy score with default hyperparameters: {0:0.4f}'.format(acc))

Model accuracy score with default hyperparameters: 0.9685
```

```
In [ ]: from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt

cm = confusion_matrix(ytest, svc_pred, labels=svc.classes_)
cm_disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=svc.classes_)

cm_disp.plot()
plt.title("SVM Model - Confusion Matrix", fontsize=16)
plt.show()
```



Part2:

1. From the above Confusion Matrix we can see that high number of Class 1 is predicted as Class 0 from the model. What is your reasoning behind this situation?

The default parameters for an SVC model do not perform as well as fine-tuned parameters.

1. What can be done in order to resolve this issue?

Tinker with the parameters to determine which would yield the best results for predicting.

SVM with high margin

Q7 Create SVM Model on the training set, and evaluate

Note:

1. If we analyze our dataset using `df.describe()` function, we will see that there are many outliers in the dataset.
2. So, we need to increase our margin with HIGH `C` values so that the SVM model get better generalization

Task:

1. Now, Create a SVM Model with rbf kernel and `C=100`
2. Fit the model with the train data
3. Predict the values with the test data
4. Calculate the model accuracy on test data
5. Plot Confusion Matrix from the true & predicted test data (Make font size 16)

A7 Replace ??? with code in the code cell below

```
In [ ]: # instantiate classifier with rbf kernel and C=100
svc = SVC(C=100.0, kernel='rbf')

# fit classifier to training set
svc.fit(xtrain, ytrain)

# make predictions on test set
rbf_pred = svc.predict(xtest)

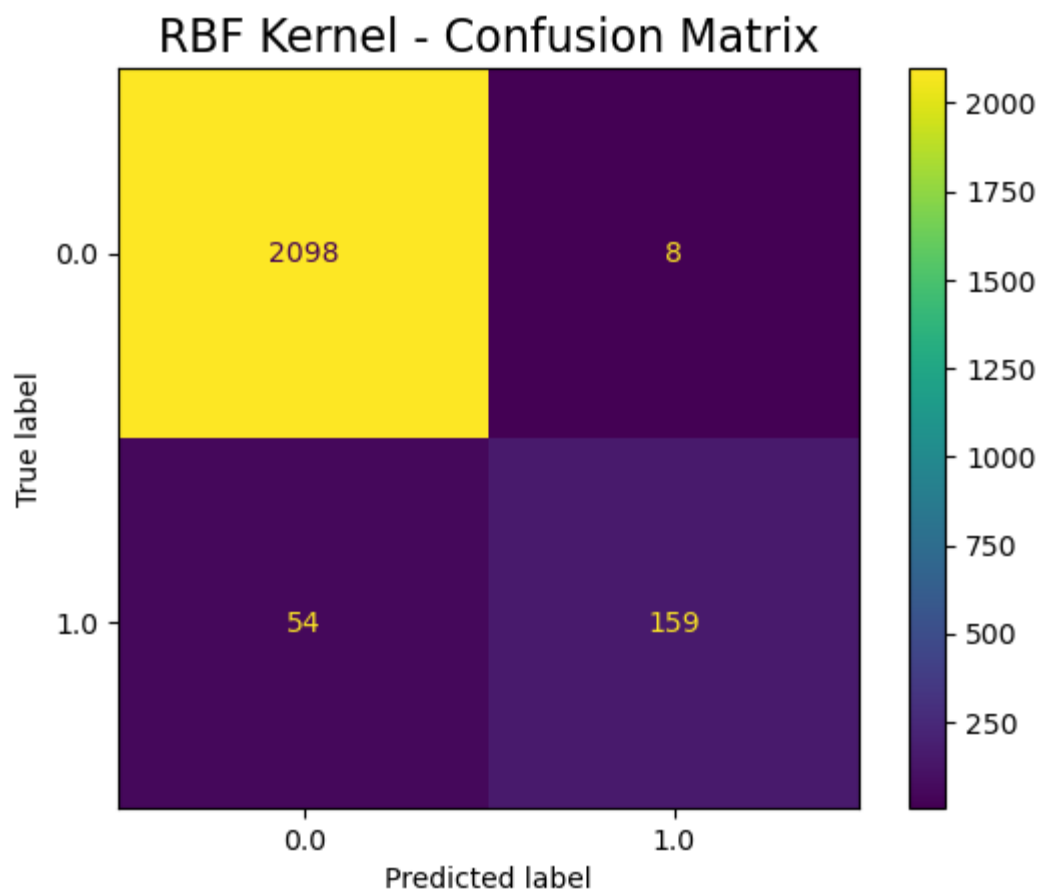
# compute and print accuracy score
acc = accuracy_score(ytest, rbf_pred)
print('Model accuracy score with rbf kernel and C=100.0 : {0:0.4f}'.format(acc))

Model accuracy score with rbf kernel and C=100.0 : 0.9733
```

```
In [ ]: from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt

cm = confusion_matrix(ytest, rbf_pred, labels=svc.classes_)
cm_disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=svc.classes_)

cm_disp.plot()
plt.title("RBF Kernel - Confusion Matrix", fontsize=16)
plt.show()
```



SVM with linear kernel

Q8 Create SVM Model on the training set, and evaluate

Task:

1. Now, Create a SVM Model with linear kernel and $C=1.0$
2. Fit the model with the train data
3. Predict the values with the test data
4. Calculate the model accuracy on test data
5. Plot Confusion Matrix from the true & predicted test data (Make font size 16)

A8 Replace ??? with code in the code cell below

```
In [ ]: # instantiate classifier with linear kernel and C=1.0
linear_svc = SVC(C=1.0, kernel='linear')

# fit classifier to training set
linear_svc.fit(xtrain, ytrain)

# make predictions on test set
linear_pred = linear_svc.predict(xtest)

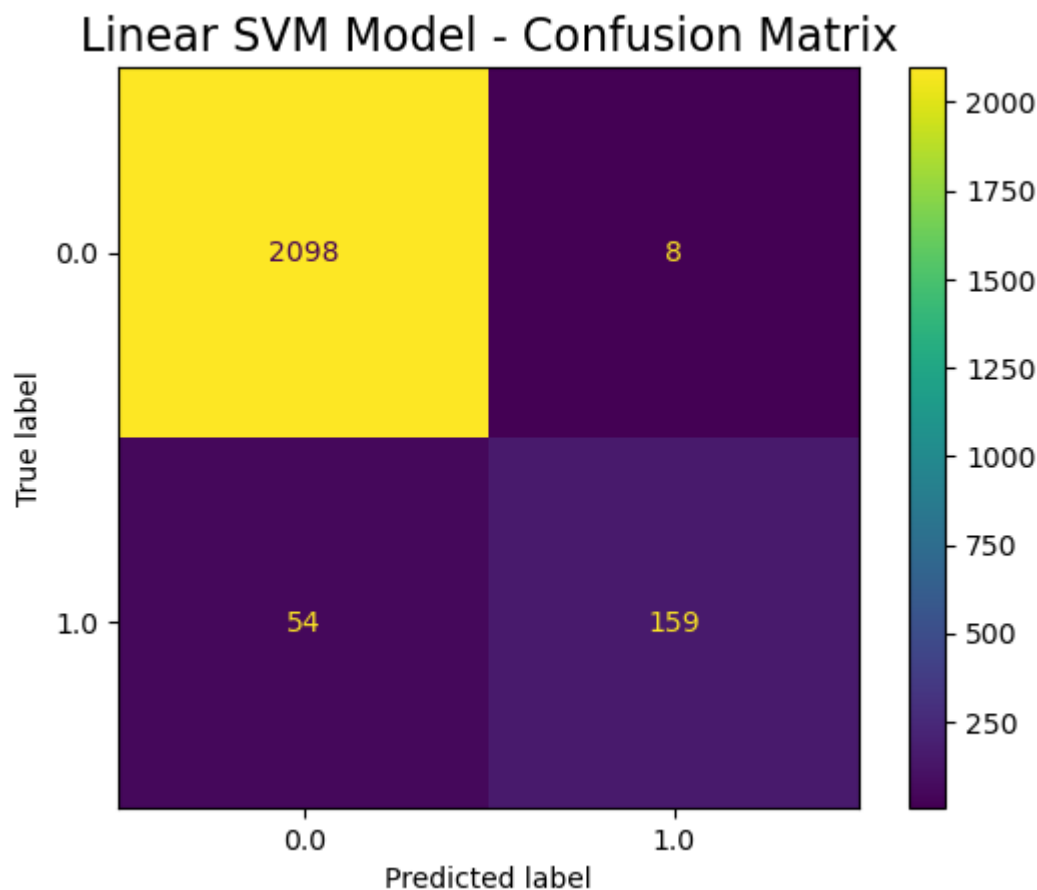
# compute and print accuracy score
lin_acc = accuracy_score(ytest, linear_pred)
print('Model accuracy score with linear kernel and C=1.0 : {0:0.4f}'.format(lin_acc))

Model accuracy score with linear kernel and C=1.0 : 0.9741
```

```
In [ ]: from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt

cm = confusion_matrix(ytest, rbf_pred, labels=svc.classes_)
cm_disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=svc.classes_)

cm_disp.plot()
plt.title("Linear SVM Model - Confusion Matrix", fontsize=16)
plt.show()
```



Q9 Create a Grid Search for finetuning the value of `C` in SVM Model on the training set ,

Task:

1. Now, Create a SVM Model with linear kernel and evaluate the model for different values of `C`. Use '`C`': `[0.01, 0.1, 5, 10, 100]`
2. Use the [sklearn GridSearchCV](#) method for finetuning the linear SVM .
3. Use 3 fold of Cross Validation
4. Use accuracy as the scoring technique
5. Use `clf.cv_results_` & `clf.best_params_` for getting the fine-tuned results from the Cross Validation run.
6. Now, Plot the Confusion Matrix for test data, using the best value of `C` we found from our finetune.

Note: The Grid Search may take couple of minutes. Please wait until the cell compiles

A9 Replace ??? with code in the code cell below

```
In [ ]: from sklearn.model_selection import GridSearchCV
        from sklearn import svm

        # Select the optimal C parameter by cross-validation
        tuned_parameters = {'C':[0.01, 0.1, 5, 10, 100]}
        svc = svm.SVC(kernel='linear')
        clf = GridSearchCV(svc, tuned_parameters, scoring='accuracy', cv=3)
```

```
In [ ]: clf.fit(xtrain, ytrain)
```

```
Out[ ]: ▶ GridSearchCV
        ▶ estimator: SVC
          ▶ SVC
```

```
In [ ]: clf.cv_results_
```

```
Out[ ]: {'mean_fit_time': array([ 0.26187881,  0.80521607, 16.79767434, 28.79816222,
                                137.95772982]),
        'std_fit_time': array([ 0.0424564 ,  0.04078544,  2.14925627,  7.25246291, 14.8237
                                9017]),
        'mean_score_time': array([0.02810756, 0.02735662, 0.02347398, 0.02578712, 0.022703
                                25]),
        'std_score_time': array([0.00225142, 0.00437647, 0.00208902, 0.00194161, 0.0012843
                                1]),
        'param_C': masked_array(data=[0.01, 0.1, 5, 10, 100],
                                mask=[False, False, False, False, False],
                                fill_value='?',
                                dtype=object),
        'params': [{'C': 0.01}, {'C': 0.1}, {'C': 5}, {'C': 10}, {'C': 100}],
        'split0_test_score': array([0.97713546, 0.98101812, 0.98101812, 0.98101812, 0.9827
                                4374]),
        'split1_test_score': array([0.97670406, 0.97929249, 0.98188093, 0.98188093, 0.9801
                                5531]),
        'split2_test_score': array([0.97411562, 0.97799827, 0.97886109, 0.97842968, 0.9788
                                6109]),
        'mean_test_score': array([0.97598504, 0.9794363 , 0.98058671, 0.98044291, 0.980586
                                71]),
        'std_test_score': array([0.00133357, 0.00123703, 0.00127003, 0.0014665 , 0.0016141
                                7]),
        'rank_test_score': array([5, 4, 1, 3, 1])}
```


```
In [ ]: clf.best_params_
```

```
Out[ ]: {'C': 5}
```

```
In [ ]: best_model = SVC(C=5, kernel='linear')
best_model.fit(xtrain,ytrain)
best_pred = best_model.predict(xtest)

cm = confusion_matrix(ytest, best_pred, labels=best_model.classes_)
cm_disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=best_model.cla

cm_disp.plot()
plt.title("SVM Model - Confusion Matrix", fontsize=16)
plt.show()
```



We can see that after using the Best Value of C , we have less amount of false positive in our test data prediction.