

HW 1

Canvas assignment [link](#) for submission and due date.

Objective:

- familiarize with the use of basic functions in python packages
- familiarize with the use of python notebook: compile all cells to show results

Assessment:

- successful run of the Markdown and Python notebook cells.
- succesfull submission on git [link](#) by the due date.

Keep the following in mind for **all** notebooks you develop:

1. Structure your notebook. Use headings with meaningful levels in Markdown cells, and explain the questions each piece of code is to answer or the reason it is there.
2. Make sure your notebook can always be rerun from top to bottom.
3. **DO NOT** erase notebook cells provided.

Setup

This section loads the relevant Python modules and does any configuration needed for the notebook to work.

Lets import python packages we will use in this homework:

- [numpy](#) - scientific computing package
- [pandas](#) - python data analysis package
- [seaborn](#) - statistical data visualization package

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
```

Introduction to Pandas

In this chapter you will use pandas commands

Q1: Read the data file using Pandas. **Note** When we run your experiment to test for correctness, we assume that the `day.csv` is in the `../data/` folder relative to your `HW1.ipynb`.

1. **Download the Capital Bike Share data set** from <https://archive.ics.uci.edu/ml/datasets/bike+sharing+dataset>. Click 'Data Folder', download the zip file, and extract the `day.csv` file.

- get used to downloading data files and saving them to correct hierarchy.
- big part of project and source versioning practice

2. Read the Data File: use Pandas `read_csv()` function to read the file into `bikes` dataframe.

- Our data file does not have column headers, so we need to specify the names.

A1 Replace the ? mark with your answer in the python cell below.

```
In [ ]: bikes = pd.read_csv('day.csv')
```

Q2: Use `head` to show the first few rows of the table:

- brief preview is a safety check you are exploring the correct data frame

A2 Replace the ? mark with your answer

```
In [ ]: bikes.head()
```

```
Out[ ]:   instant  dteday  season  yr  mnth  holiday  weekday  workingday  weathersit  temp
0         1  2011-01-01      1   0     1         0         6           0           2  0.344167  0.
1         2  2011-01-02      1   0     1         0         0           0           2  0.363478  0.
2         3  2011-01-03      1   0     1         0         1           1           1  0.196364  0.
3         4  2011-01-04      1   0     1         0         2           1           1  0.200000  0.
4         5  2011-01-05      1   0     1         0         3           1           1  0.226957  0.
```

Q3: Use `info` to show a description of the columns, along with the shape and memory use of the data frame:

- `.info()` or `.head()` can be called in the same cell as data load
- we separate them out in this notebook so that we can discuss them in the markdown cells, but we can combine them in the future.

A3 Replace the ? mark with your answer in the python cell below

```
In [ ]: bikes.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   instant    731 non-null    int64  
 1   dteday      731 non-null    object  
 2   season      731 non-null    int64  
 3   yr          731 non-null    int64  
 4   mnth        731 non-null    int64  
 5   holiday     731 non-null    int64  
 6   weekday     731 non-null    int64  
 7   workingday  731 non-null    int64  
 8   weathersit   731 non-null    int64  
 9   temp        731 non-null    float64 
10   atemp       731 non-null    float64 
11   hum         731 non-null    float64 
12   windspeed   731 non-null    float64 
13   casual      731 non-null    int64  
14   registered  731 non-null    int64  
15   cnt         731 non-null    int64  
dtypes: float64(4), int64(11), object(1)
memory usage: 91.5+ KB

```

Q4: Pandas provide a very useful function for exploring statistical properties of dataframe, and allow us to see data composition for numerical columns. Use pandas build-in function and show statistical information for columns.

A4: Replace the ? mark with your answer in the python cell below

In []: `bikes.describe()`

Out[]:

	instant	season	yr	mnth	holiday	weekday	workingday	weath
count	731.000000	731.000000	731.000000	731.000000	731.000000	731.000000	731.000000	731.00
mean	366.000000	2.496580	0.500684	6.519836	0.028728	2.997264	0.683995	1.35
std	211.165812	1.110807	0.500342	3.451913	0.167155	2.004787	0.465233	0.54
min	1.000000	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	1.00
25%	183.500000	2.000000	0.000000	4.000000	0.000000	1.000000	0.000000	1.00
50%	366.000000	3.000000	1.000000	7.000000	0.000000	3.000000	1.000000	1.00
75%	548.500000	3.000000	1.000000	10.000000	0.000000	5.000000	1.000000	2.00
max	731.000000	4.000000	1.000000	12.000000	1.000000	6.000000	1.000000	3.00

Q5: Use pandas functions for filtering the dataframe rows based on some column values. Find out number of rows where the value of the **temp** column is more than the average value of **temp** for the dataset.

Steps:

1. Find out the mean value of the temp column
2. Filter the rows where temp is greater than the mean value
3. Get the number of rows in the filtered dataframe.

A5: Replace the ? mark with your answer in python cell below

```
In [ ]: mean_value= bikes['temp'].mean()
        filtered_dataframe= bikes[bikes['temp'] > mean_value]
        row_count= len(filtered_dataframe)

        print(row_count)
```

367

Numpy

We now use Numpy for doing some mathematical calculation on the dataset.

Q6: Now, use numpy for working with below tasks

1. At first convert the dataframe into a numpy array
2. Print the shape of the numpy n-dimensional array
3. select and print rows from 100 to 105

A6: Replace the ? mark with your answer

```
In [ ]: num_array= bikes.to_numpy() # conver the dataframe
        print(num_array.shape) #print the shape
        print(num_array[99:105,:]) # print rows from 100 to 105

(731, 16)
[[100 '2011-04-10' 2 0 4 0 0 0 2 0.426667 0.426737 0.8575 0.146767 1188
  1707 2895]
 [101 '2011-04-11' 2 0 4 0 1 1 2 0.595652 0.565217 0.716956 0.324474 855
  2493 3348]
 [102 '2011-04-12' 2 0 4 0 2 1 2 0.5025 0.493054 0.739167 0.274879 257
  1777 2034]
 [103 '2011-04-13' 2 0 4 0 3 1 2 0.4125 0.417283 0.819167 0.250617 209
  1953 2162]
 [104 '2011-04-14' 2 0 4 0 4 1 1 0.4675 0.462742 0.540417 0.1107 529 2738
  3267]
 [105 '2011-04-15' 2 0 4 1 5 0 1 0.446667 0.441913 0.67125 0.226375 642
  2484 3126]]
```

Q7: Lets put it all together

1. Create a new numpy array selecting column number 10 - 13.
2. Sort the numpy array in ascending order based on the 2nd column of our new numpy array.
3. Print first 5 rows of the sorted numpy array

A7: Replace the ? mark with your answer in the cell below

```
In [ ]: new_array= bikes.iloc[:,9:13].values

sorted_array= new_array[new_array[:, 1].argsort()]

print(sorted_array[:5,:])

[[0.0591304 0.0790696 0.4          0.17197   ]
 [0.0965217 0.0988391 0.436522   0.2466    ]
 [0.1275    0.101658  0.464583   0.409212   ]
 [0.138333  0.116175  0.434167   0.36195    ]
 [0.0973913 0.11793   0.491739   0.15833    ]]
```

Seaborn Plotting 1


Seaborn package is used to plot the data. For every question in this chapter use the **bikes dataframe** for answering the questions.

Q8: Make a bar plot showing the mean number of riders (y-axis) per weekday (x-axis) using seaborn `catplot` method.

A8: Replace the ? mark with your answer

```
In [ ]: mean_riders = sns.catplot(data=bikes, x='weekday', y='cnt', kind='bar', estimator='
mean_riders.set_axis_labels('Day of the Week','Mean Number of Riders')

Out[ ]: <seaborn.axisgrid.FacetGrid at 0x13c2f878950>
```



Analysis

The X-axis labels are not very helpful as day 0 is not clear. This is a question about how the data is *coded*. We'll talk more about data encoding next week. Unfortunately, the data documentation doesn't actually say how weekdays are coded! But we can infer from the data in this case: first data point is January 1, 2011, which was a Saturday, coded as weekday 6; it then resets to 0 for the next day, and starts counting up. Often, we will not be able to infer the data encoding from the data itself - we need to consult the codebook or data set description. We got lucky this time. But looking at the data can help us make sense of the codebook.

Lesson here is to always look at your data.

Q9: Turn these weekday numbers into a *categorical* variable so Pandas knows how to label them. Hint: use `pandas.Categorical.from_codes()`.

A9: Replace the ? mark with your answer in the python cell below

```
In [ ]: codes = pd.CategoricalDtype(['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Fr  
bikes['day_names'] = pd.Categorical.from_codes(codes=bikes['weekday'], dtype=codes)  
bikes.head()
```

```
Out[ ]:
```

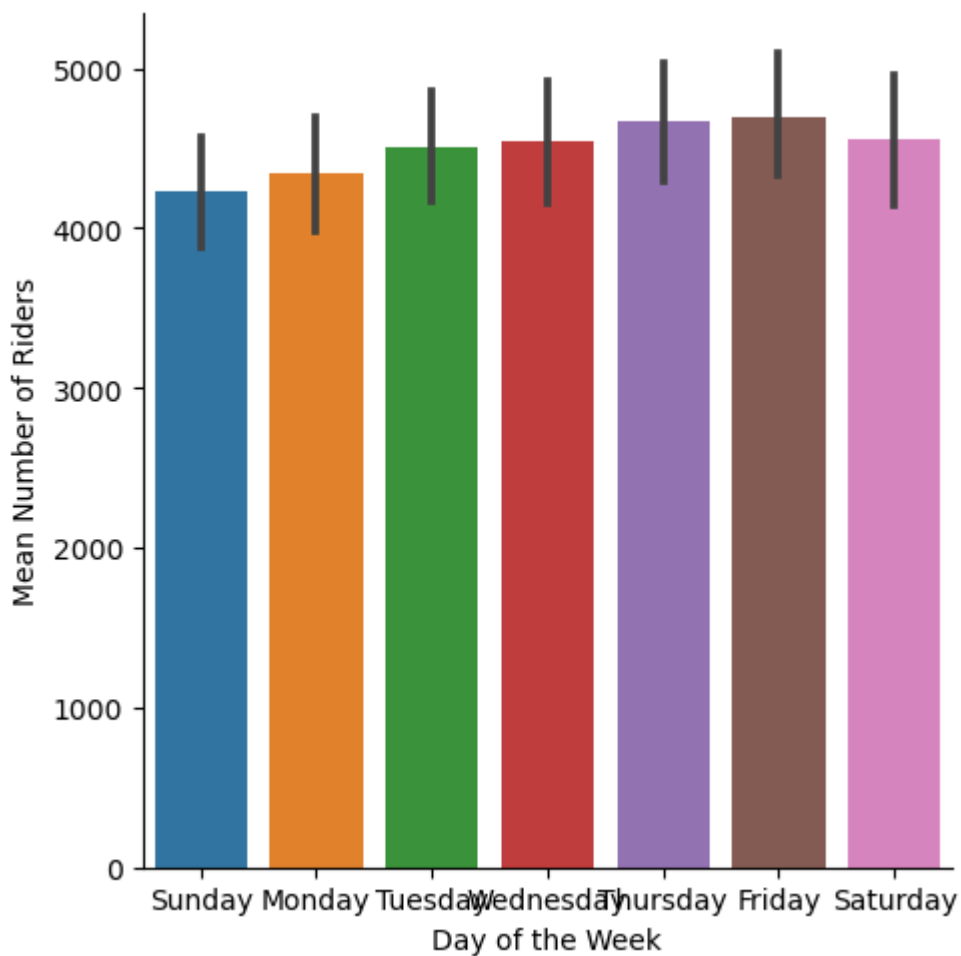
	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp
0	1	2011-01-01	1	0	1	0	6	0	2	0.344167
1	2	2011-01-02	1	0	1	0	0	0	2	0.363478
2	3	2011-01-03	1	0	1	0	1	1	1	0.196364
3	4	2011-01-04	1	0	1	0	2	1	1	0.200000
4	5	2011-01-05	1	0	1	0	3	1	1	0.226957

A10: Plot new data using seaborn `catplot` , where data=bikes, x-axis is `day_names` and y-axis is `cnt`

A10: Replace the ? mark with your answer in the python cell below

```
In [ ]: mean_riders = sns.catplot(data=bikes, x='day_names', y='cnt', kind='bar', estimator=
mean_riders.set_axis_labels('Day of the Week', 'Mean Number of Riders'))
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x13c2f9b2b90>
```



You have now now plotted the average rides per day.

Note: When we do not tell `catplot` what to do with multiple points for the same value (in this case the weekday name), it computes the mean and a bootstrapped 95% confidence interval.

Seaborn Polotting 2: View Data over Time

Lets explore how did rides-per-day change over the course of the data set?


- This kind of data - a sequence of data points associated with times - is called a *time series*.
- This data set gives us an `instant` column that records the data number since the start of the data set

Q11: Use `seaborn.lineplot()` where `data=bikes`, `x-axis` is `instant` and `y-axis` is `cnt` value.

A11: Replace the ? mark with your answer in the python cell below

```
In [ ]: sns.lineplot(data=bikes, x='instant', y='cnt')
```

```
Out[ ]: <Axes: xlabel='instant', ylabel='cnt'>
```



Lets view this graph for actual times on x-axis. The `dteday` column records the date. We can transform `dteday` column to the actual date comlumn using `pandas.to_datetime()` method on the column:


```
In [ ]: bikes['dt'] = pd.to_datetime(bikes['dteday'])
```

Q12: Now create a plot using `seaborn.lineplot()` where data=bikes, x-axis is `dt` and y-axis is `cnt` value.

A12: Replace the ? mark with your answer

```
In [ ]: sns.lineplot(data=bikes, x='dt', y='cnt')
```


```
Out[ ]: <Axes: xlabel='dt', ylabel='cnt'>
```



Next, plot the *weekly* rides by resampling. Right now, our `bikes` data is indexed by row number in the CSV file. We can change its index to another column, such as our `dt` column with the date, which then lets us do things like resample by week:

```
In [ ]: bikes.set_index('dt')['cnt'].resample('1W').sum().plot()
```

```
Out[ ]: <Axes: xlabel='dt'>
```



What that code did, in one line, is:

1. Set the data frame's index to `dt` (`bikes.set_index('dt')`), returning a new DF
2. Select the count column (`['cnt']`), returning a series
3. Resample the series by week (`.resample('1W')`)
4. Combine measurements within each sample by summing them (`.sum()`)
5. Plotting the results using Pandas' defaults (`.plot()`)

Pandas default plotting functions are useful for quick plots to see what's in a data frame or series. They often are difficult to use to turn in to publication-ready charts.

Q13: Save the modified dataframe in csv format using pandas `to_csv()` function. Give the file name as **day_output**

A13: Replace the ? mark with your answer

```
In [ ]: bikes.to_csv('day_output.csv', index=True)
        # I opted to include row numbers, index=False to exclude
```

Submission Instructions

1. Run all cells in `HW1.ipynb` and make sure there are no errors
2. Print `HW1.ipynb` to pdf file
3. Upload `HW1.ipynb` , `HW1.pdf` and `day_output.csv` files to your course assigned git repo e.g: <https://git.txstate.edu/ML/netid> before the deadline.